

CoActo: CoActive Neural Network Inference Offloading with Fine-grained and Concurrent Execution

Mobisys'24

Kyungmin Bin
Seoul National University
kmbin@snu.ac.kr

Jongseok Park
Seoul National University
cakeng@snu.ac.kr

Chanjeong Park
Seoul National University
cjpark99@snu.ac.kr

Seyeon Kim
University of Colorado Boulder
seyeon.kim@colorado.edu

Kyunghan Lee
Seoul National University
kyunghanlee@snu.ac.kr

1 Introduction


Mobile can provide high-quality services that are comparable to those of human experts now



user interactions is important


Core target: Reduce the inference latency

1 Introduction



On-device Inference
(For more **complex models**, unable
to **achieve low latency**)

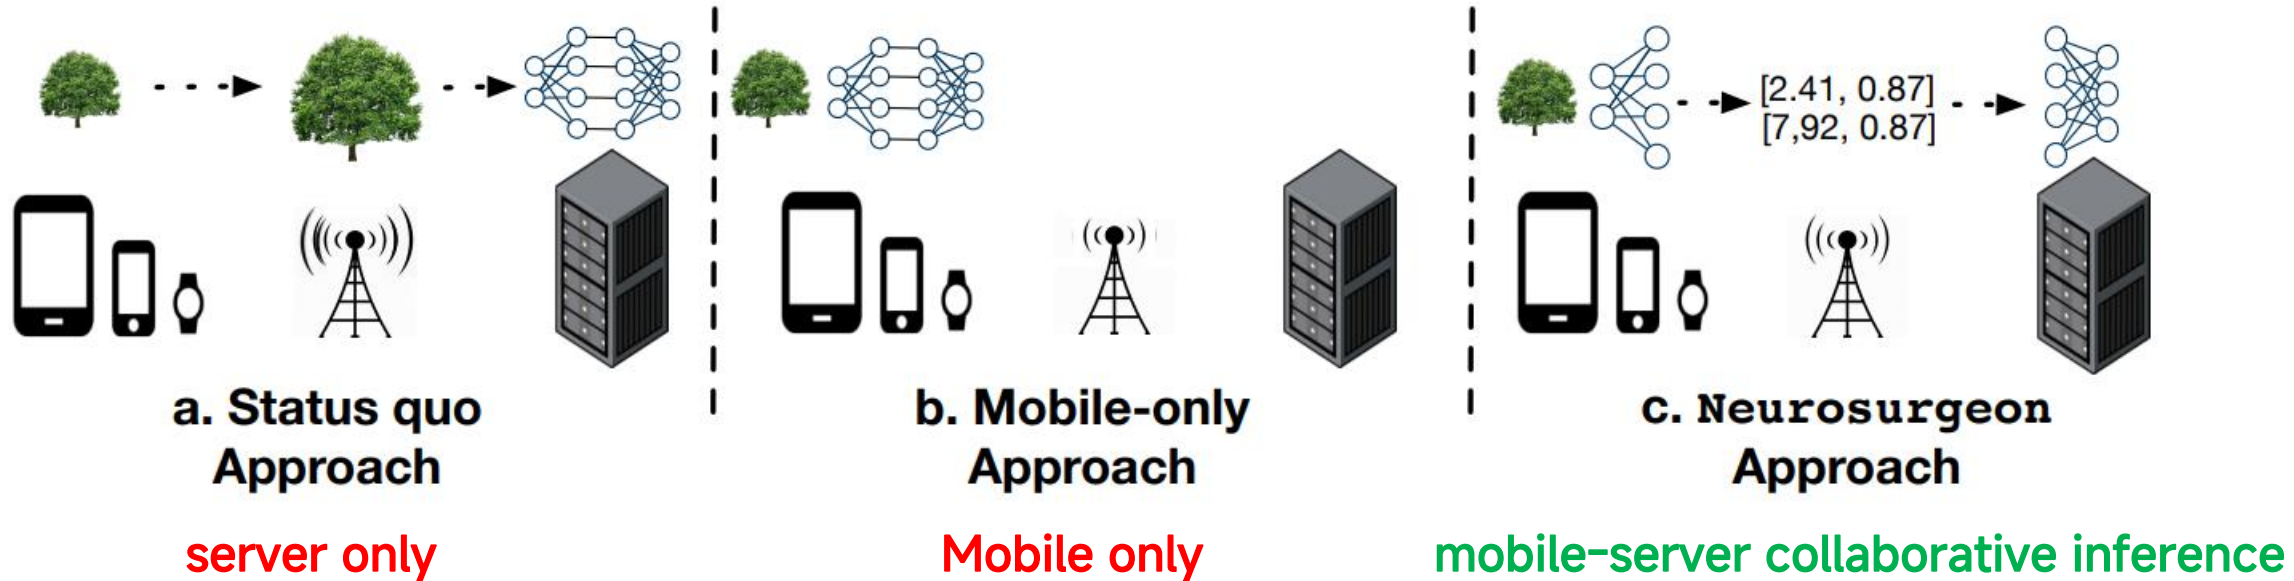
VS



Offload to server
(With **more computing power** to
support **more complex tasks**)

1 Introduction


Offload all the model to the cloud is still not the best choice



- mobile inference
- server inference
- transmission cost
- optimal partitioning
- best scheduling scheme

1 Introduction

For DNN offloading, this not only includes the **partitioning and scheduling of the workload**, but also the **modeling of the workload, execution algorithm, dynamic load-balancing, possibility of multi-tenant execution, and many more.**



**C. Neurosurgeon
Approach**

mobile-server collaborative inference

- mobile inference
- server inference
- transmission cost
- **optimal partitioning**
- best scheduling scheme

A more fundamental question:

How should a DNN execution system be designed for efficient mobile DNN offloading?

1 Introduction

A more fundamental question:

How should a DNN execution system be designed for efficient mobile DNN offloading?

To realize efficient DNN inferences in a mobile-server offloading environment

Two key properties

- a fine-grained expression of DNNs
- flexibility of the system resource utilization

1 Introduction

A more fundamental question:

How should a DNN execution system be designed for efficient mobile DNN offloading?

To realize efficient DNN inferences in a mobile-server offloading environment

Two key properties

- a fine-grained expression of DNNs
 - DNNs are often expressed in **units of layers**.
 - too **large**
 - not supply enough parallelism to efficiently utilize all available resources.
- flexibility of the system resource utilization

1 Introduction

A more fundamental question:

How should a DNN execution system be designed for efficient mobile DNN offloading?

To realize efficient DNN inferences in a mobile-server offloading environment

Two key properties

- a fine-grained expression of DNNs
- flexibility of the system resource utilization
 - **dynamic nature** of DNN offloading
 - the system must be **flexible** enough to support **dynamic changes**
 - changes in available computation resources
 - network conditions
 - competing inference offloads

2 Background and Motivation

two representative approaches in collaborative inference:




Figure 2: Illustrations of two collaborative inference approaches, (a) split computing and (b) fused-layer offloading.

- Split computing
 - splits the DNN model into two submodels **at the layer level**
 - the key is the **split point**
 - many studies try to find the the optimal split point
 - **this sequential execution cannot make full use of the available computing resources**

2 Background and Motivation

two representative approaches in collaborative inference:




Figure 2: Illustrations of two collaborative inference approaches, (a) split computing and (b) fused-layer offloading.

- Fused-layer (FL) offloading
 - **fuse** multiple layers by exploiting the spatial locality of layers
 - several submodels **with zero data dependencies**
 - executed without any **synchronization** to the computation of other submodels,

2 Background and Motivation

two representative approaches in collaborative inference:

- Fused-layer (FL) offloading




Fig. 2. A simple DNN computation with the FL technique and homogeneously intercepted fused layer's size.

2 Background and Motivation

two representative approaches in collaborative inference:




Figure 2: Illustrations of two collaborative inference approaches, (a) split computing and (b) fused-layer offloading.

- Fused-layer (FL) offloading

- **fuse** multiple layers by exploiting the spatial locality of layers
- several submodels **with zero data dependencies**
- executed **without any synchronization**
- **suffers from limited scalability and high computation overhead**

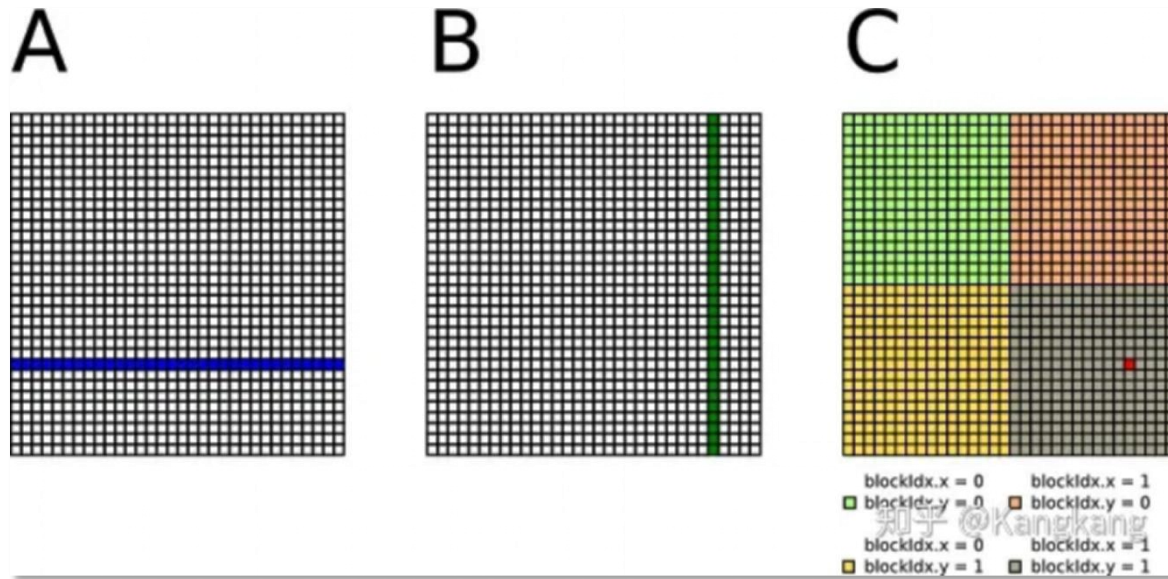
Layers are simply too **large** of a unit for a **DNN offloading environment**



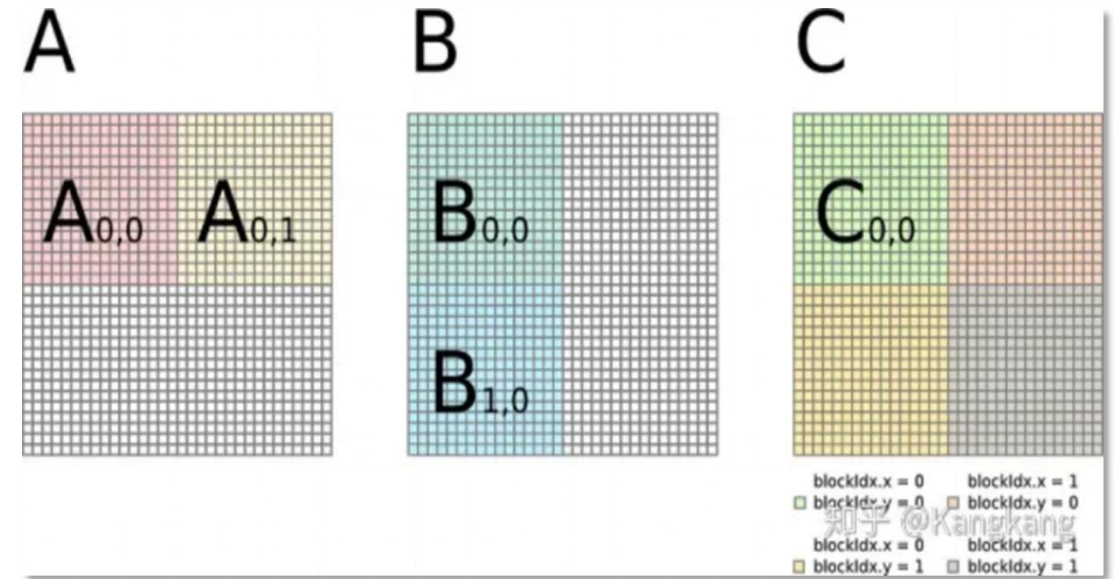
Tiling for Collaborative Inference

2 Background and Motivation

Tiling for Collaborative Inference



native matrix multiplication





use tiling technique

2 Background and Motivation

Tiling for Collaborative Inference

$$\begin{bmatrix} A_{0,0} & A_{0,1} \end{bmatrix} \begin{bmatrix} B_{0,0} \\ B_{1,0} \end{bmatrix} = \begin{bmatrix} C_{0,0} \end{bmatrix}$$

increase the parallelism
of the operation



tile sizes is flexible

tiles to be the ideal unit for fine-grained expression of DNN computation

2 Background and Motivation

Design Philosophy

Traditional approaches for collaborative inference primarily focus on **model splitting**

Not only to **distribute the workload** but also to ensure that runtime execution system components work in unison to **make the best use of the available resources** under the **dynamic environments** that exist during DNN offloading

2 Background and Motivation

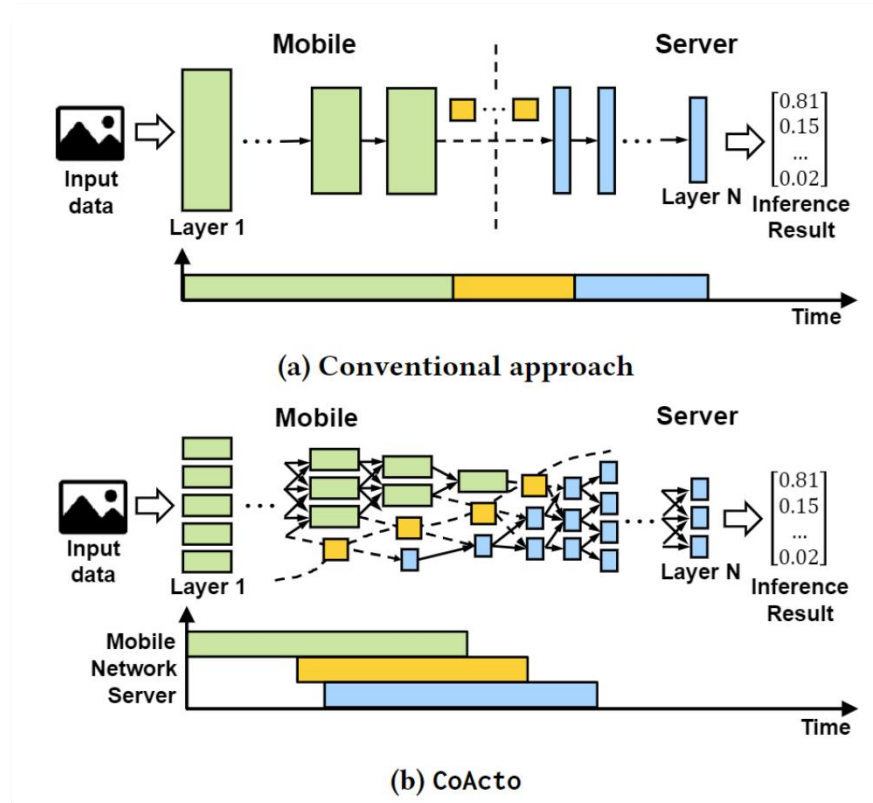
Two design philosophy

- Fine-grained DNN expression
 - **Smaller workload size** allows faster unit processing times
 - suitable for **dynamic scheduling** of parallel resources

2 Background and Motivation

Two design philosophy

- Fine-grained DNN expression
- Concurrency of runtime resources
 - **Concurrent**, rather than sequential, use of these resources is necessary to maximize parallel resource utilization.



2 Background and Motivation

Three challenges

1. Tile-based expression

- expresses an arbitrary layer-wise computation graph as a tile-wise computation graph poses many challenges
 - determining the efficient tile dimensions and size for the given environment
 - automatically parsing
 - generating the independent data dependency flow graph between the tiles

2 Background and Motivation

Three challenges

1. Tile-based expression
2. Concurrent execution system
 - Tensorflow or PyTorch execute at layer level
 - tiling restricts the concurrency only to the intra-layer level
 - designing a concurrent execution system that enables overlapping the computation and communications of tiles is challenging

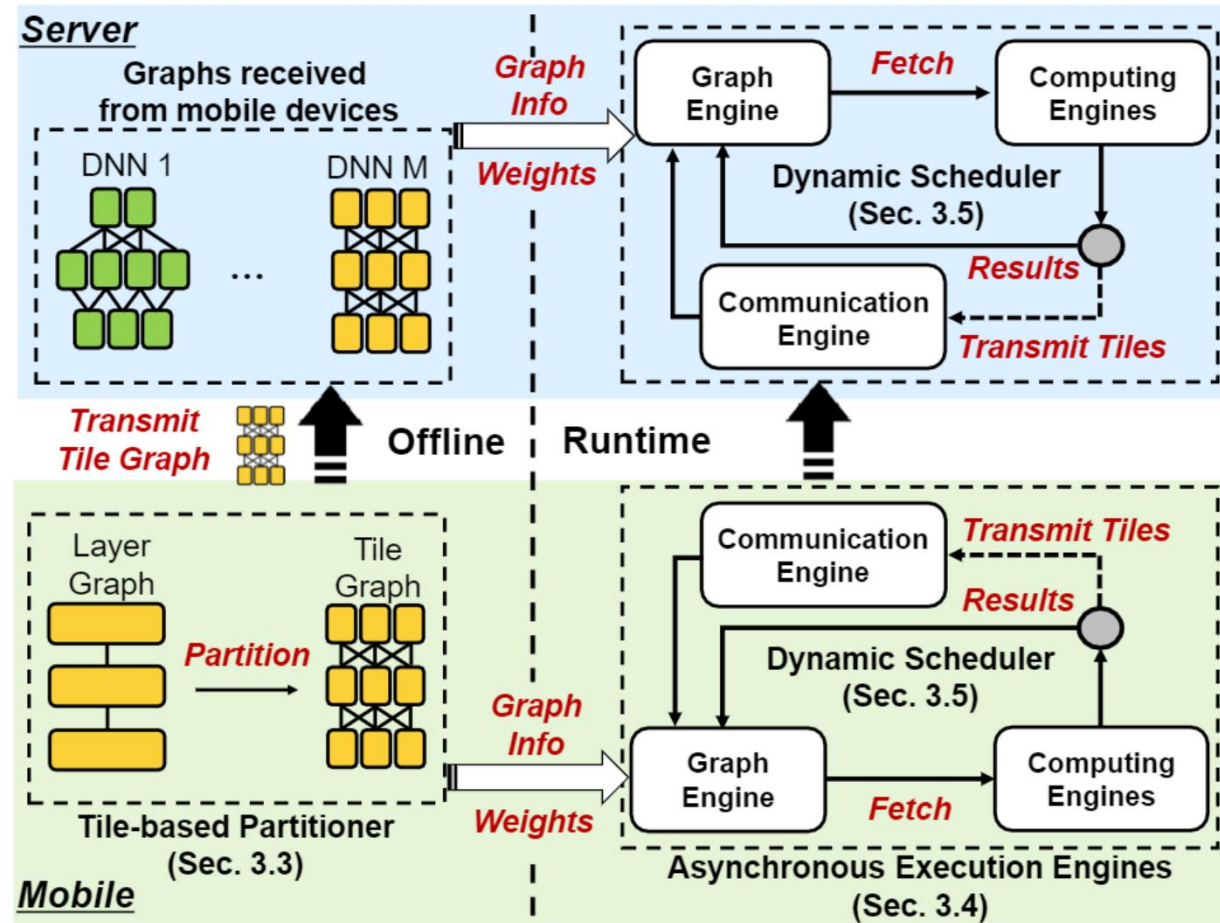
2 Background and Motivation

Three challenges

1. Tile-based expression
2. Concurrent execution system
3. **Dynamic scheduling of tiles**
 - **balancing** the model executions between the **mobile, network, and server resources** of the given environment
 - **dynamic adaptation** and **balancing of complex fine-grained DNN** between the concurrently operating resources is also challenging

3 CoActo Design

Overview



Three challenges

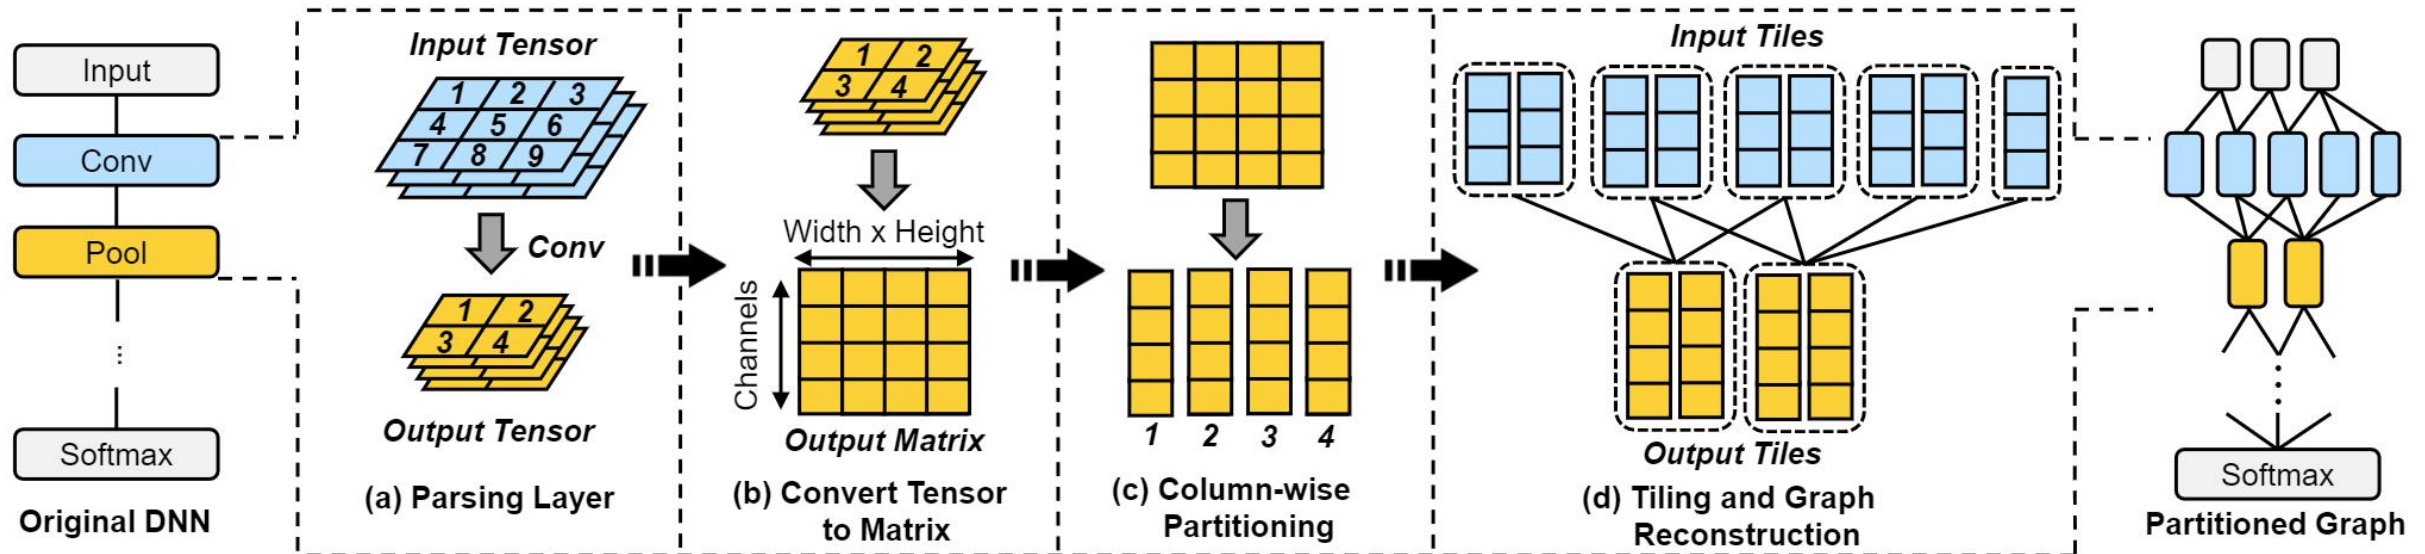
1. Tile-based expression
2. Concurrent execution system
3. Dynamic scheduling of tiles

Three components

- Tile-based Partitioner (TP)
- Asynchronous Execution Engines (AEEs)
- Dynamic Scheduler (DS)

3 CoActo Design

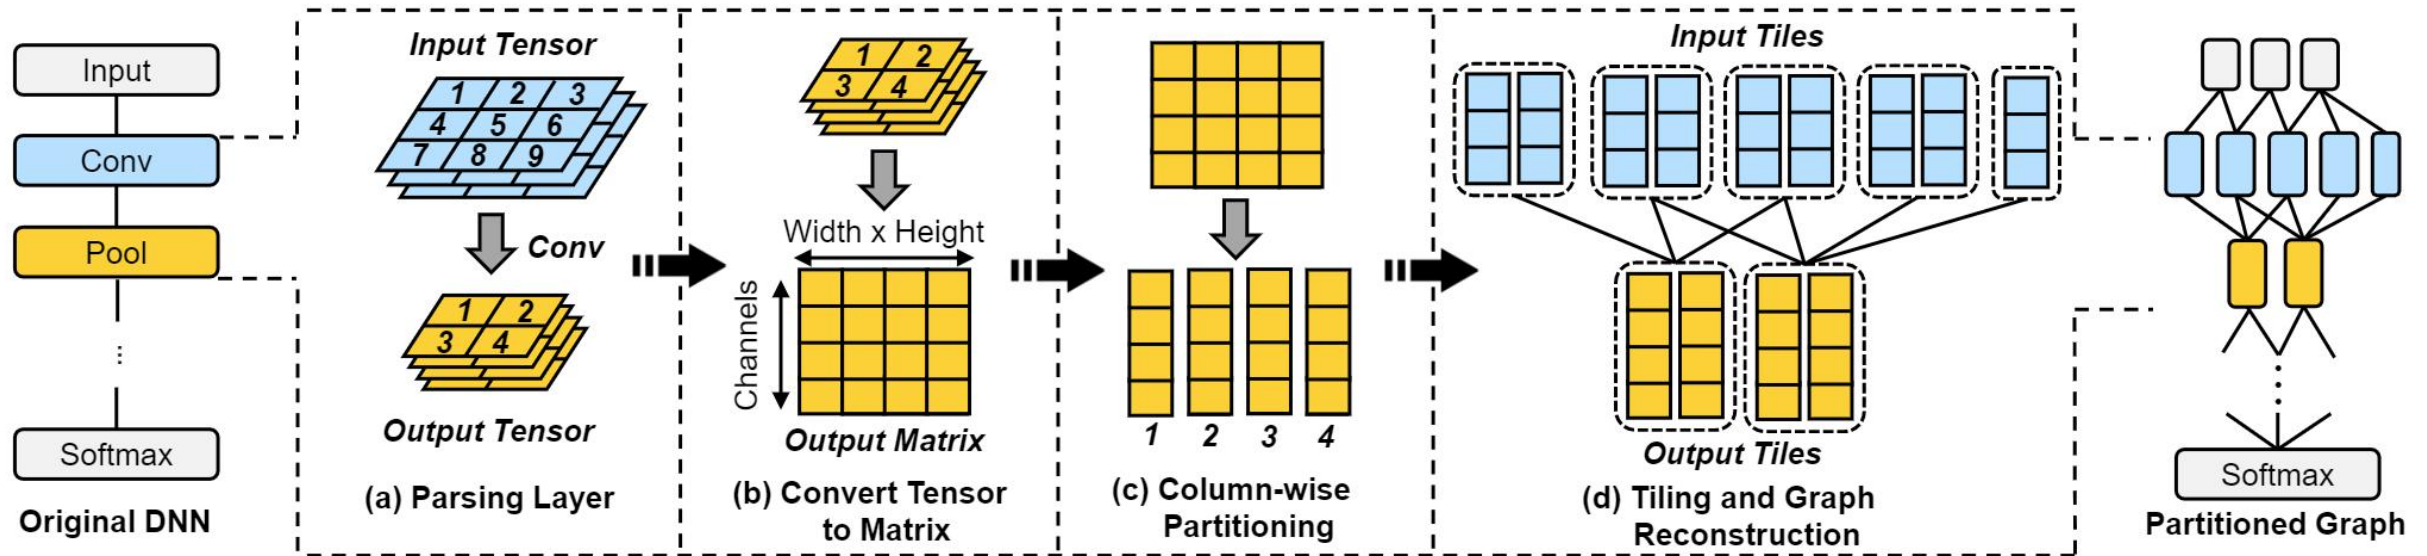
Tile-based Partitioner (TP)



1. analyze the size of each layer's output tensor
2. TP flattens the output tensors into a matrix
3. the matrices are decomposed into column-wise vectors
4. tiles are created by merging the partitioned columns (**determined by the tile number which is predefined and changes by the environment**)

3 CoActo Design

Tile-based Partitioner (TP)



The computation graph is then reconstructed by graphing the tiles into a **directed acyclic graph (DAG)**

3 CoActo Design

Tile-based Partitioner (TP)

Tiling allows a **flexible** and **concurrent execution** in DNN inference offloading

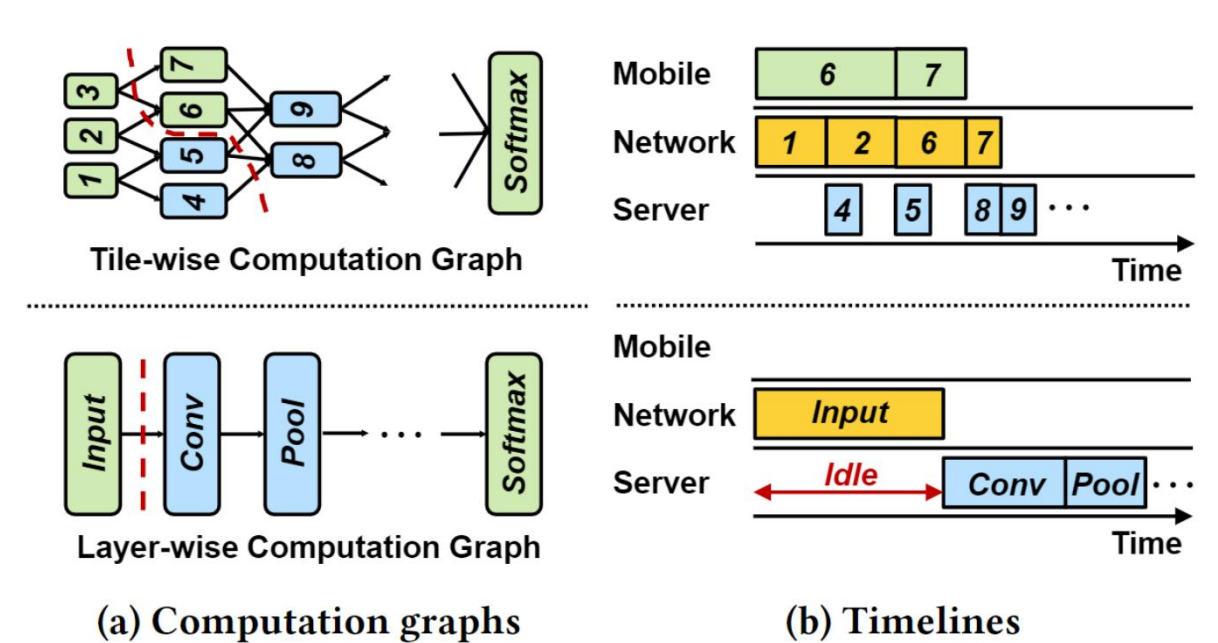
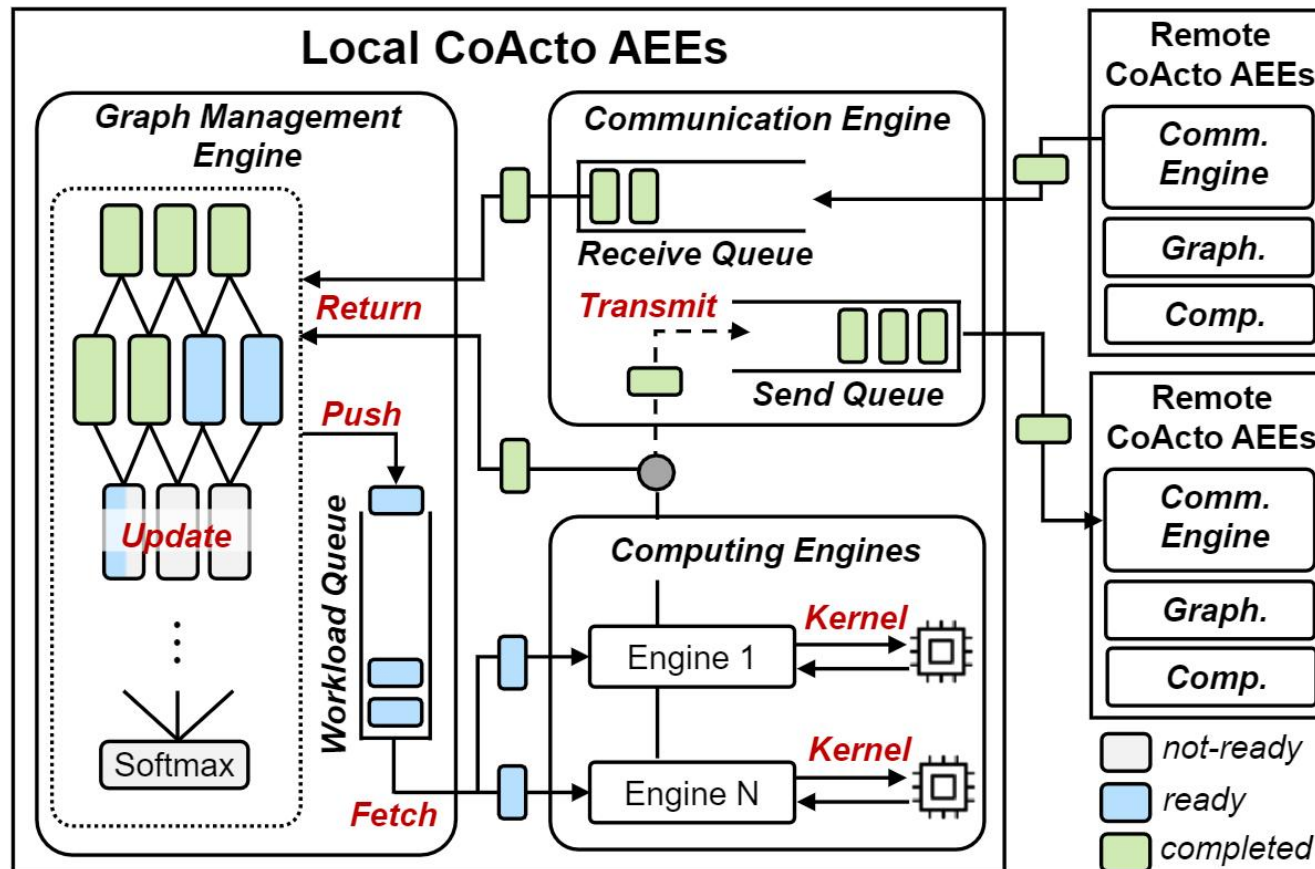


Figure 5: Examples of the collaborative inference with the tiles-wise computation graph (top) and the layer-wise computation graph (bottom).

3 CoActo Design

Asynchronous Execution Engines (AEEs)



Three types of execution engines

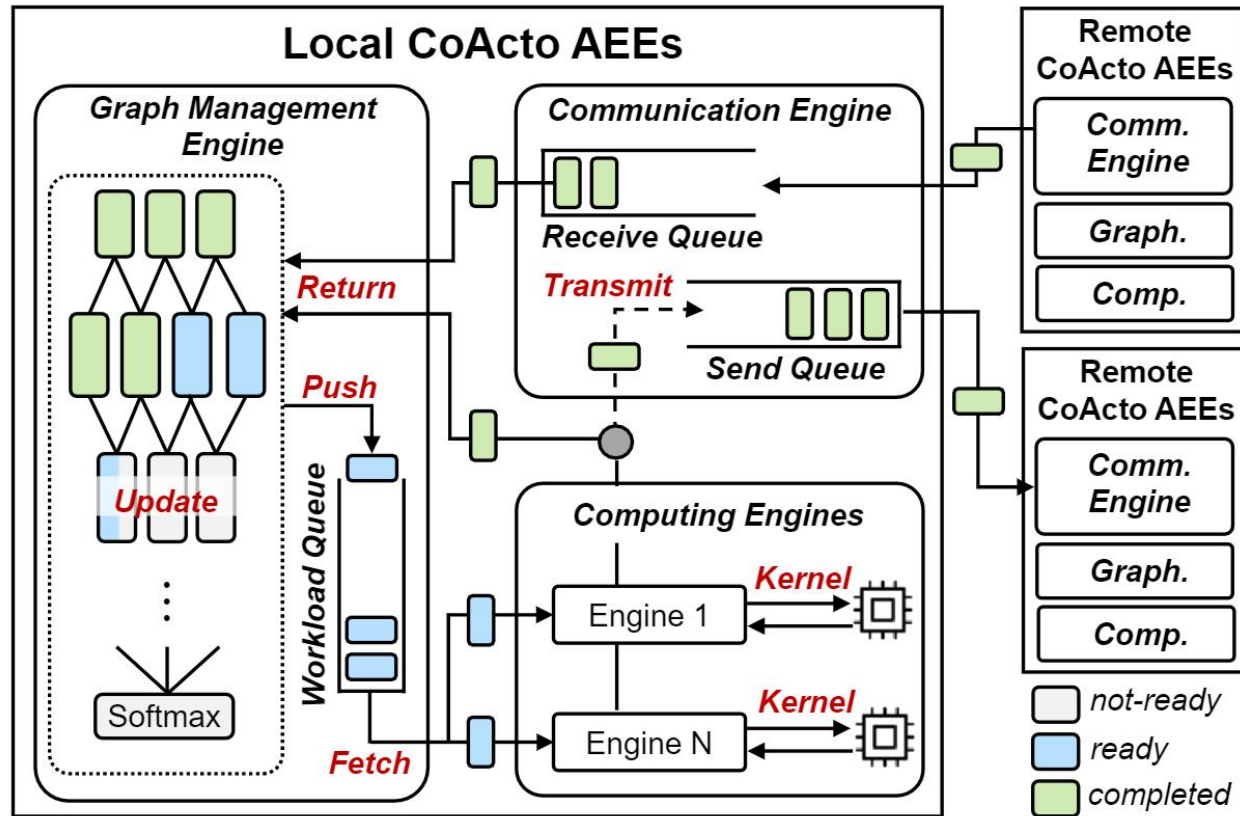
- Graph Management Engine
- Computing Engine
- Communication Engine

asynchronously and independently
operate without waiting for each other

Figure 6: The overall execution workflows of Asynchronous Execution Engines.

3 CoActo Design

Asynchronous Execution Engines (AEEs)




Graph Management Engine

- managing the computation graph **separately** in each parallel computing engine requires **frequent synchronization**
- use **Graph Management Engine**
 - contains the **entire** computation graph information and its state

Figure 6: The overall execution workflows of Asynchronous Execution Engines.

3 CoActo Design

Asynchronous Execution Engines (AEEs)



Graph Management Engine

Three state of node(tile)

- completed
- ready
- not-ready

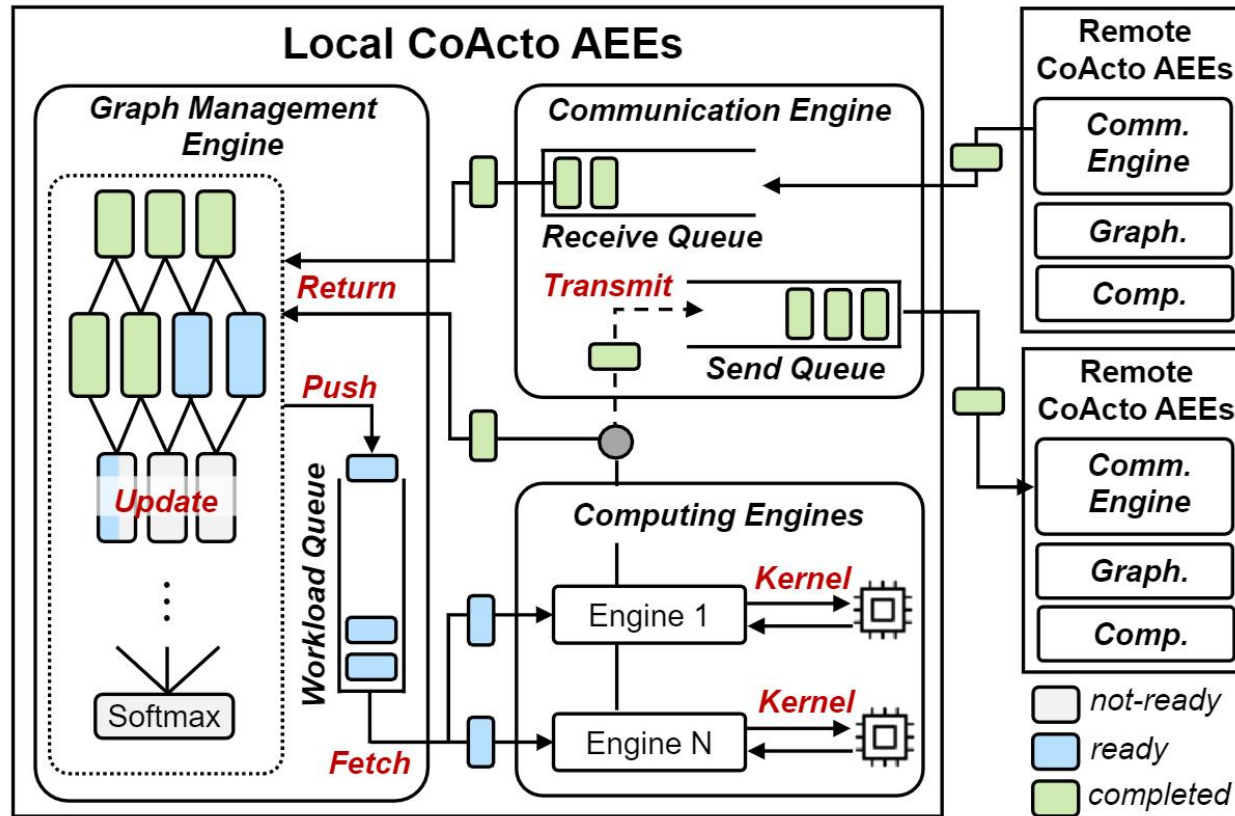
Workload queue

- pushes any child node that becomes ready to the workload queue

Figure 6: The overall execution workflows of Asynchronous Execution Engines.

3 CoActo Design

Asynchronous Execution Engines (AEEs)



Computing & Communication Engines

All computing and communication engines **operate concurrently and asynchronously** without any **synchronization**

Figure 6: The overall execution workflows of Asynchronous Execution Engines.

3 *CoActo* Design


Dynamic Scheduler (DS)

With the **tile-based** partitioning technique, scheduling the tiles is interpreted as **a complex DAG scheduling problem** that is a well-known **NP-complete** problem

- Static DAG scheduling approaches is not a good choice for **dynamic environments**
 - **unexpected network interference**
 - **an extremely burst request**
- **A dynamic offloading decision algorithm is better**

3 *CoActo* Design

Dynamic Scheduler (DS)



(a) Sample DAG

Task model

We define the partitioned computation graph as a **DAG** $G = \langle V, E \rangle$

V : *nodes(tiles)*


E : *data dependency*

v point with no child is called *vertex*

Each node v_i has a **computation cost (FLOPs)** of c_i , and output tile **data size** of d_i

3 CoActo Design

Dynamic Scheduler (DS)



(a) Sample DAG

Goal of the task

o_i denotes the offloading decision of a node v_i

$CT(v_i)$ represents the completion time of a node v_i

find the optimal offloading policy $O = \{o_1, \dots, o_N\}$ that **minimizes the maximum completion time of the exit nodes v_{exit}** .

$$\min_{o_i \in O} \max CT(v_{exit}) \quad (1)$$

3 *CoActo* Design

Dynamic Scheduler (DS)

Basic idea

- first send **all the input nodes** to the server (sharing input data is cheap)
- **dynamically** decide to compute and send the outputs of the subsequent nodes on runtime

3 CoActo Design

Dynamic Scheduler (DS)

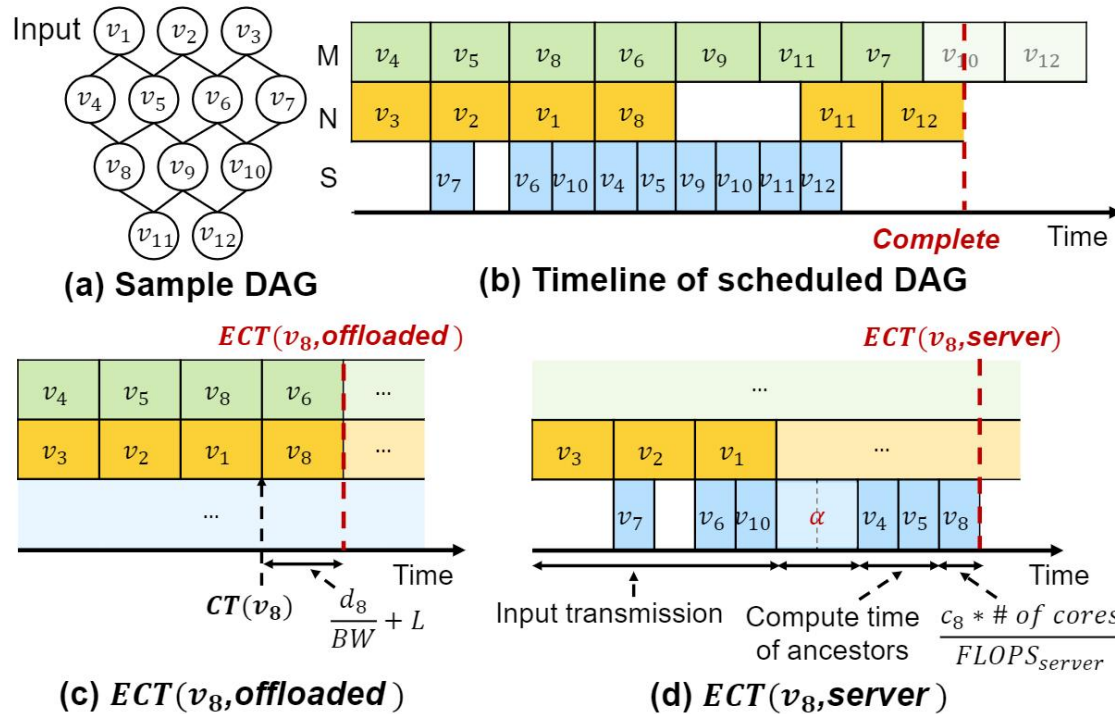


Figure 7: An example of dynamic offloading decision with (a) a sample DAG. It is performed by calculating $ECT(v_i, \text{offloaded})$ and $ECT(v_i, \text{server})$.

Dynamic offloading decision

- Mobile and server compute with the **largest diameter** in between to minimize **duplicated computation**
 - e.g. (b) mobile trans v_3 while executing v_4 , server execute v_7 after v_3 arrives

3 CoActo Design

Dynamic Scheduler (DS)




Figure 7: An example of dynamic offloading decision with (a) a sample DAG. It is performed by calculating $ECT(v_i, \text{offloaded})$ and $ECT(v_i, \text{server})$.

Dynamic offloading decision

- using **a greedy approach** in the mobile device
 - decision is made using the **estimated completion times**

$$ECT(v_i, \text{offloaded}) < ECT(v_i, \text{server})$$

e.g v8

enables the maximal utilization of the powerful server resources

3 CoActo Design

Dynamic Scheduler (DS)




Figure 7: An example of dynamic offloading decision with (a) a sample DAG. It is performed by calculating $ECT(v_i, \text{offloaded})$ and $ECT(v_i, \text{server})$.

Estimating the completion times

- $ECT(v_8, \text{offloaded})$
 - the completion time $CT(v_8)$
 - the queuing latency
 - obtained by **total data size** of the nodes in the send queue of and the **profiled bandwidth**.
- the transmission time

3 CoActo Design

Dynamic Scheduler (DS)




Figure 7: An example of dynamic offloading decision with (a) a sample DAG. It is performed by calculating $ECT(v_i, \text{offloaded})$ and $ECT(v_i, \text{server})$.

Estimating the completion times

- $ECT(v_8, \text{server})$
 - the transmission time of all input nodes
 - delayed time αk of the mobile device k by the resource contention (**dynamic**)
 - the computation time of ancestors of v_8
 - the computation time of v_8

4 Evaluation

Experimental Setup

| Platform | CPU | Memory |
|-------------------|---------------------------------|--------|
| Server | 64 Cores AMD Threadripper 3990X | 128GB |
| Jetson AGX Xavier | 8 Cores Carmel ARMv8.2 | 32GB |
| Raspberry Pi 4 | 4 Cores ARM Cortex-A72 1.8GHz | 8GB |
| Pixel 5 | 1 Core ARM Cortex A-76 2.4GHz | 8GB |
| | 1 Core ARM Cortex A-76 2.2GHz | |
| | 6 Cores ARM Cortex A-55 | |

Table 1: Specifications of the tested platforms.

4 Evaluation

Baselines

- **Cloud-only**: A status-quo approach that offloads whole DNN inference workloads to the cloud server by transmitting the input data.
- **On-device**: An approach that executes local inference on mobile platforms without offloading
- ***SPINN***: The state-of-the-art split computing in collaborative inference
- ***FL-offloading***: Fused-Layer (FL)-based collaborative inference approach

4 Evaluation

End-to-End Latency

Effectiveness in computation bottleneck

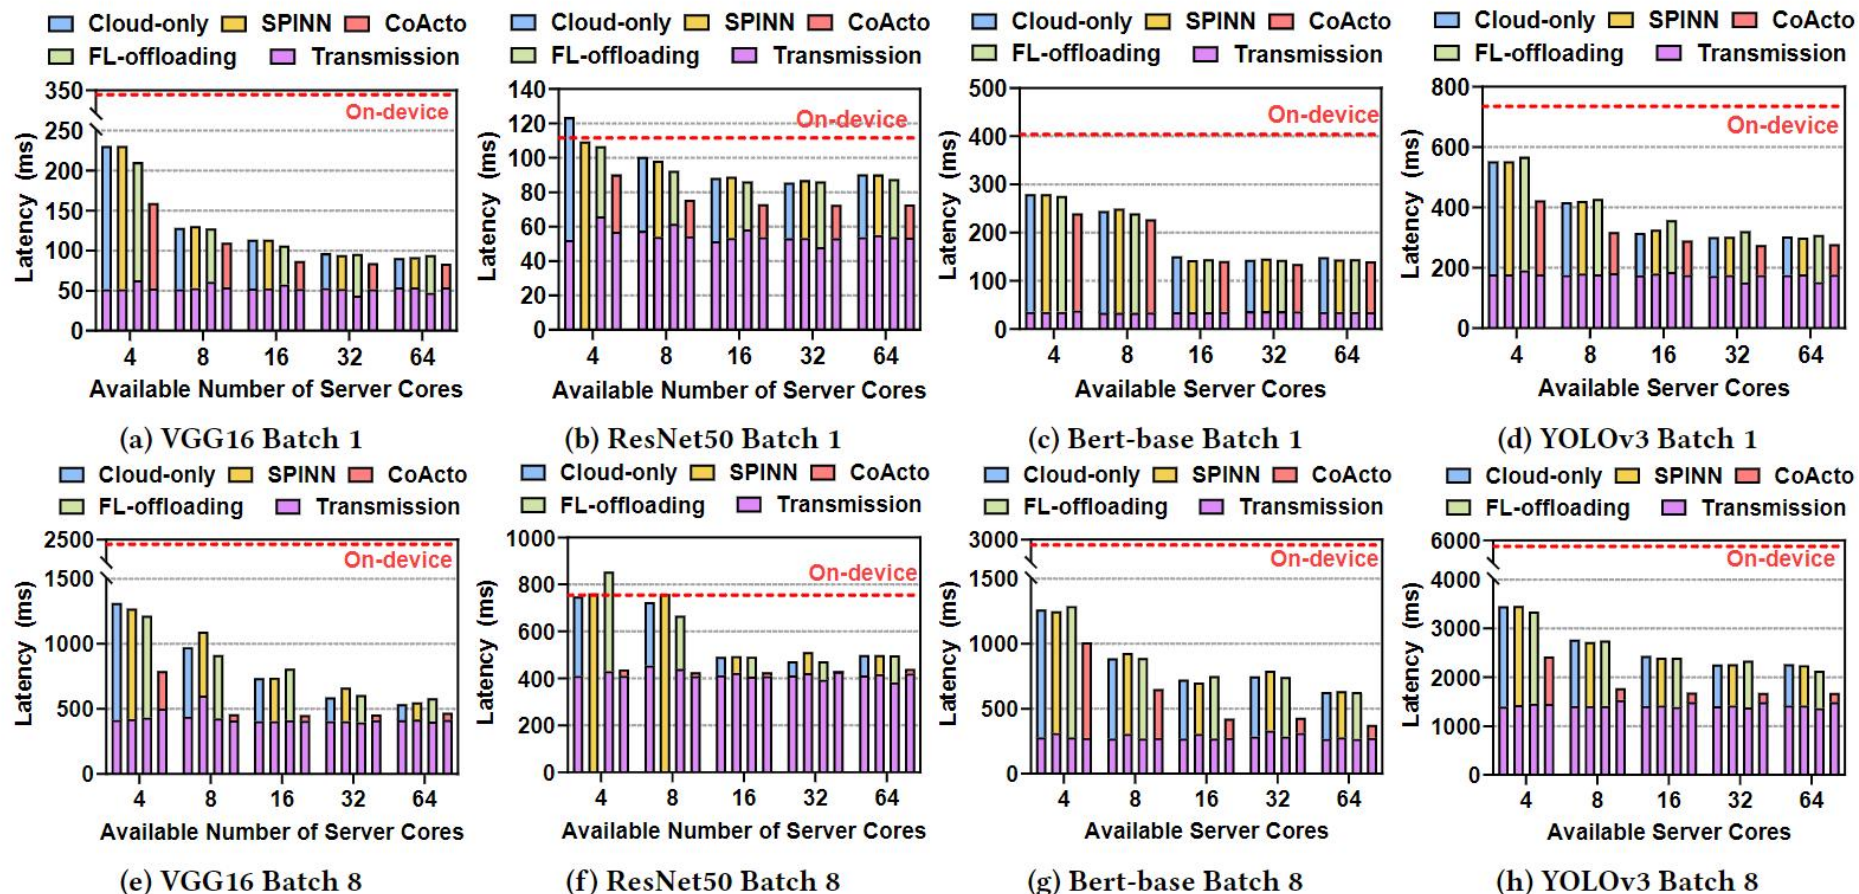


Figure 8: End-to-end latency using Jetson AGX Xavier with different number of available server cores, under a 100Mbps WiFi network.

4 Evaluation

End-to-End Latency

Effectiveness in network bottleneck

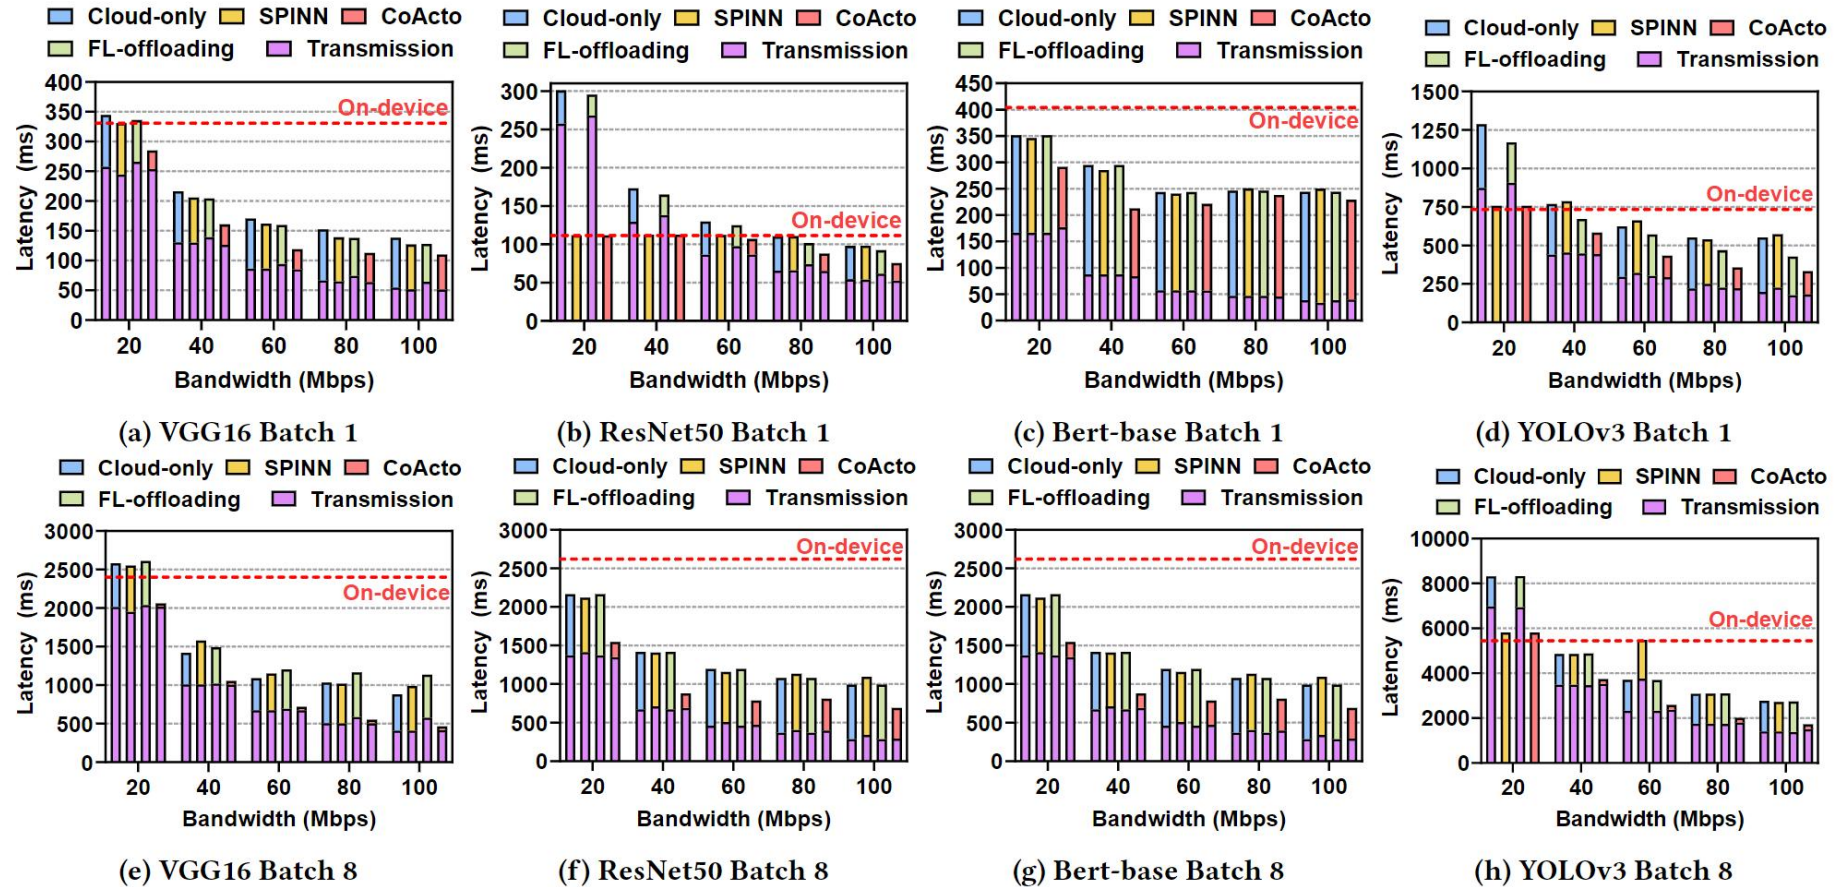


Figure 9: End-to-end latency using Jetson AGX Xavier under different network bandwidths and 8 cores available in the server.

4 Evaluation

Concurrency in Multi-tenant Scenario

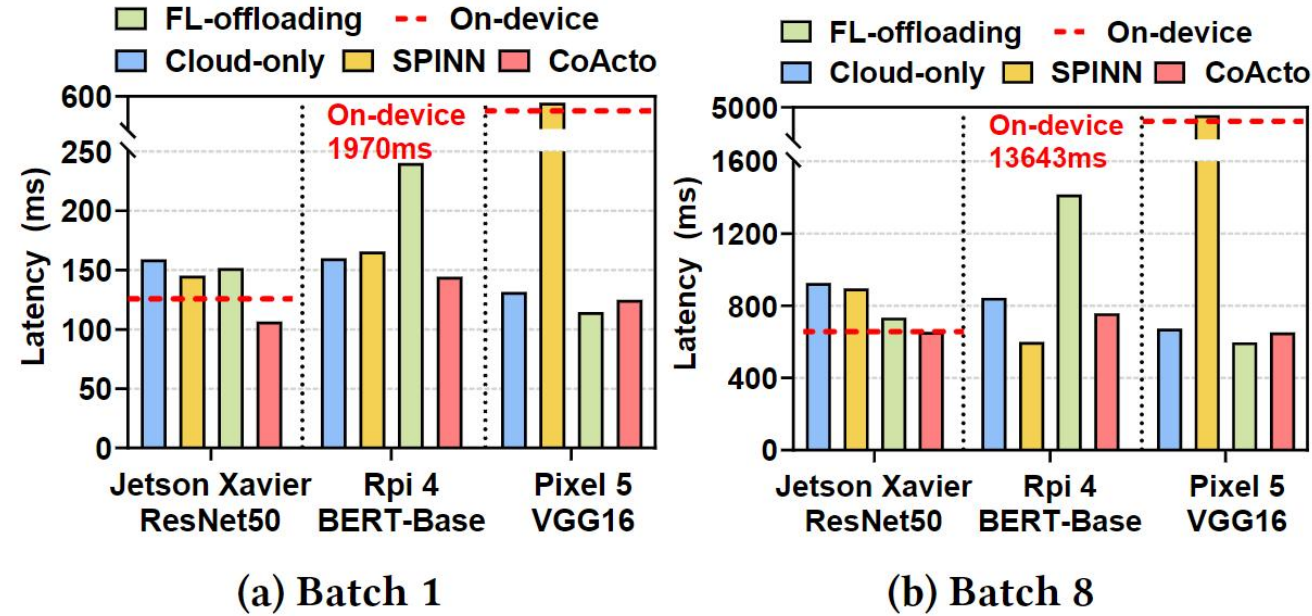


Figure 11: End-to-end latency in multi-DNN inference scenario that each device requests a distinct DNN inference query to the shared server with 100Mbps and 64 available cores settings.

4 Evaluation

Effectiveness of the Granularity

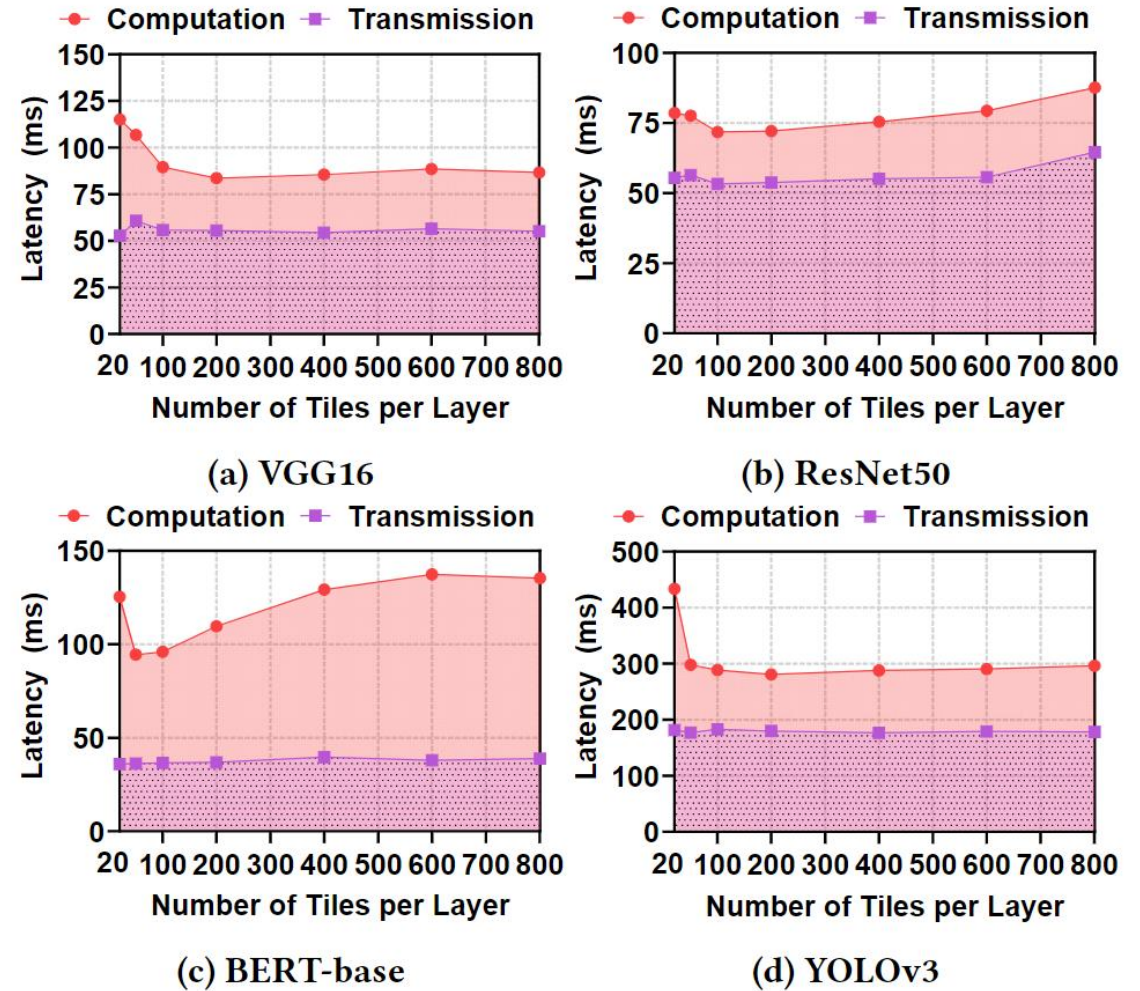


Figure 12: End-to-End latency of CoActo with different number of tiles per layer in 100MBps and 64 cores in the server. Note that the batch size is 1 for all the tested DNNs.

5 Thoughts

strengths

- optimize the mobile-server **collaborative inference** from the system view
- propose a fine-grained DNN expression tiling during execution
- support dynamic scheduling and concurrency of all the runtime resources

Weakness

- only support **CPU inference now**
 - can extend to GPU inference by implementing the tile-based computation kernels
- the system is **less effective** when the network is **bad** or when there is a **large difference** between the network transmission speed and the inference speed

Thank You.

2024.5.17