

# SpaceEvo: Hardware-Friendly Search Space Design for Efficient INT8 Inference

*ICCV'23*

Li Lyna Zhang<sup>1\*‡</sup> Xudong Wang<sup>2\*§</sup> Jiahang Xu<sup>1</sup> Quanlu Zhang<sup>1</sup> Yujing Wang<sup>3</sup>  
Yuqing Yang<sup>1</sup> Ningxin Zheng<sup>1</sup> Ting Cao<sup>1</sup> Mao Yang<sup>1</sup>

*From Microsoft Research Asia*

# Introduction

- INT8 quantization: compressing models by **reducing the number of bits** required to represent weights or activations

|      |      |      |
|------|------|------|
| 0.34 | 3.75 | 5.64 |
| 1.12 | 2.7  | -0.9 |
| -4.7 | 0.68 | 1.43 |

FP32



Quantization

|    |     |     |
|----|-----|-----|
| 64 | 134 | 217 |
| 76 | 119 | 21  |
| 3  | 81  | 99  |


INT8

Speed up

Reduce the computations  
and the inference time

# Introduction

- INT8 quantization only achieve only marginal speedup
- Traditional quantization methods focus on **minimizing accuracy** loss for a given pre-trained model, but **ignore the real-world inference efficiency**.



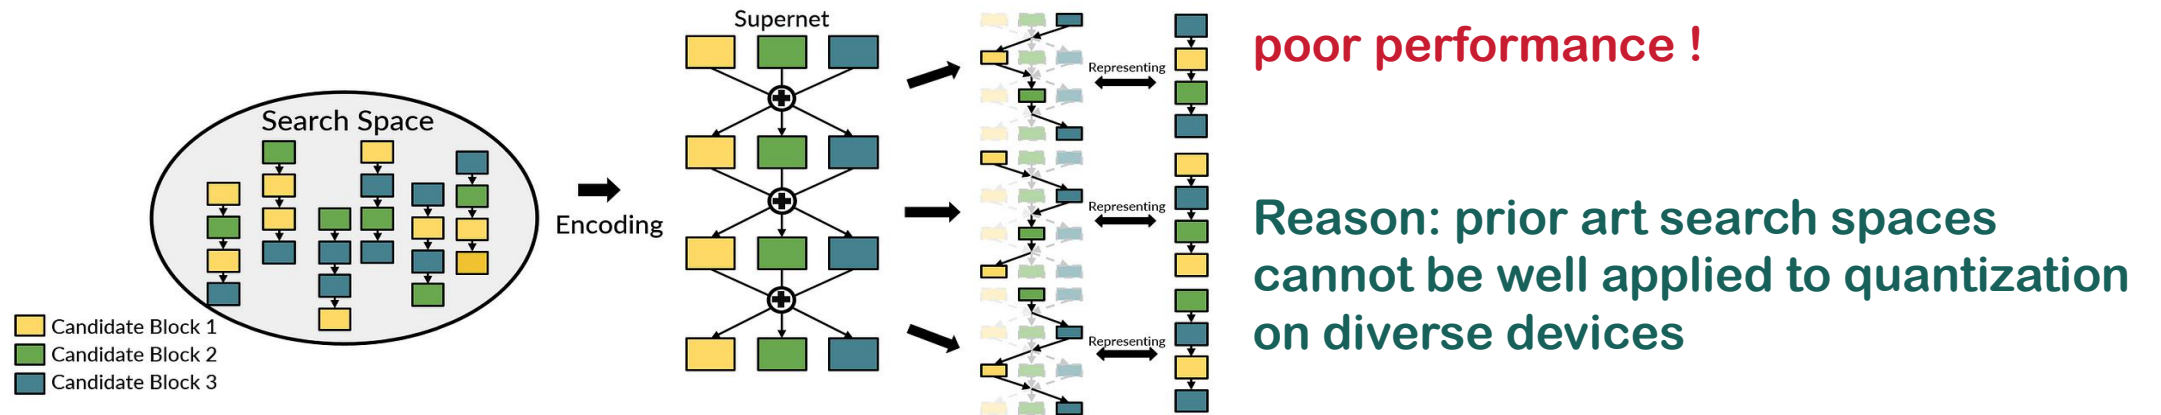
Less than 2x speedup

Still high latency !

**Difficult to deploy on latency-critical scenarios**

# Introduction

- Neural Architecture Search (NAS) is a powerful tool for automating efficient quantized model design.
- Two-stage paradigm (one shot NAS)
  - **step1**: trains a weight-shared quantized supernet
  - **step2**: uses typical search algorithms to find subnets with best quantized accuracy under different FLOPs constraints



**Current search space limit NAS to find better quantized models for edge devices.**

# Introduction


## Question:

Can we design a **hardware-friendly search space**,  
allowing NAS to discover **better models** that meet the  
low INT8 **latency and accuracy** requirements ?

# On-device Quantization Efficiency Analysis

**What factors lead to current NAS search space quantization unfriendly issue?**

**Observation 1: FP32 latency and FLOPs are not good indicators of INT8 latency**



(b) FP32 and INT8 latency on Pixel 4

A common belief is that model with lower FLOPs or FP32 latency leads to less INT8 latency


Neither of them is a good indicator of INT8 latency

# On-device Quantization Efficiency Analysis


**What factors lead to current NAS search space quantization unfriendly issue?**

**Observation 2: The choices of operators' types and configurations greatly impact the INT8 latency**


- The search space comprises a sequence of **blocks**.
- The **block type varies** and is allowed to search from **a range of hyperparameter configurations**
  - eg. kernel size, expansion ratio, channel width and depth...



residual



residual bottleneck



MBConv

# On-device Quantization Efficiency Analysis

**What factors lead to current NAS search space quantization unfriendly issue?**

**Observation 2: The choices of operators' types and configurations greatly impact the INT8 latency**

Many block type and configuration choices in current search spaces unexpectedly **slow down** the INT8 latency !

| Operator  | Intel CPU | Pixel 4 |
|-----------|-----------|---------|
| Conv      | 2.6x      | 2.5x    |
| DWConv    | 1.2x      | 2.0x    |
| SE        | 0.7x      | 1.4x    |
| Hardswish | 0.7x      | 0.5x    |
| Swish     | 0.8x      | 2.1x    |

(a) Avg. INT8 latency speedup (relative to FP32) of major operators in MobileNetV3.

Some lightweight operators become slower because of the data transformation between INT32 and INT8

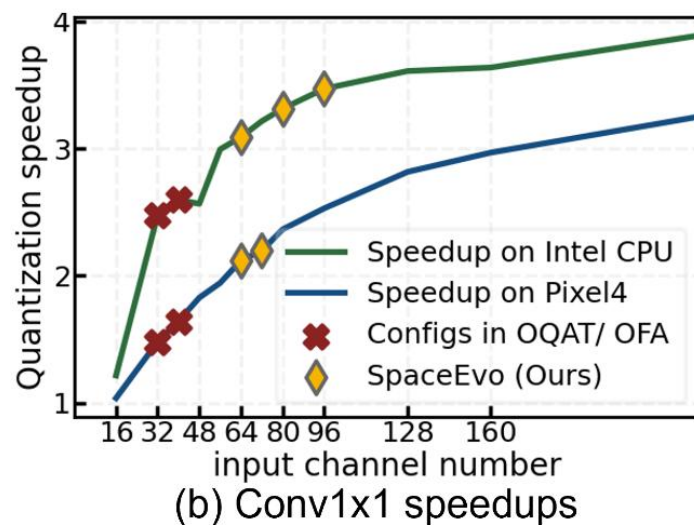


# On-device Quantization Efficiency Analysis

**What factors lead to current NAS search space quantization unfriendly issue?**

**Observation 2: The choices of operators' types and configurations greatly impact the INT8 latency**

**The configuration choices also determine the quantization efficiency**



Small channel widths in OFA search space cannot benefit well from quantization

SpaceEvo can automatically design a search space with larger channel widths for better efficiency

# On-device Quantization Efficiency Analysis

**What factors lead to current NAS search space quantization unfriendly issue?**

**Observation 3: Quantization-friendly settings are diverse and contradictory across devices**

**The quantization-friendly operators are different and can be contradictory on diverse devices.**

| Operator  | Intel CPU | Pixel 4 |
|-----------|-----------|---------|
| Conv      | 2.6x      | 2.5x    |
| DWConv    | 1.2x      | 2.0x    |
| SE        | 0.7x      | 1.4x    |
| Hardswish | 0.7x      | 0.5x    |
| Swish     | 0.8x      | 2.1x    |

Quantization speedups are highly dependent on the inference engines and hardware

(a) Avg. INT8 latency speedup (relative to FP32) of major operators in MobileNetV3.

# On-device Quantization Efficiency Analysis

**What factors lead to current NAS search space quantization unfriendly issue?**

Observation 1: FP32 latency and FLOPs are not good indicators of INT8 latency

Observation 2: The choices of operators' types and configurations greatly impact the INT8 latency

Observation 3: Quantization-friendly settings are diverse and contradictory across devices

*There is no single structure (block types in a model) that is optimal for quantization on all hardware*

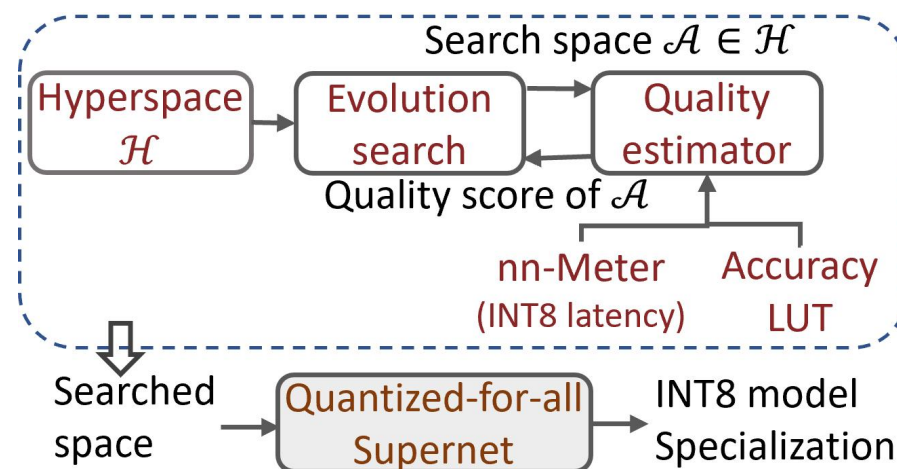
Design a *specialized quantization-friendly search space* for each hardware.

Each search space is tailored to the unique characteristics of the hardware and includes an optimal structure with elastic depths, widths, and kernel sizes.

# Methodology

## Core design concept

- Different from architecture search, where the goal is to **find the single best model from the space**
- we aim to find the **best search space** from lots of **all kinds of search spaces (hyperspace)**
  - 不同于神经网络搜索在一个搜索空间中找到最好的模型
  - 我们的目标是在一个不同搜索空间组成的更大的搜索空间 (hyperspace) 中, 找到最好的一个搜索空间作为量化模型NAS的搜索空间。



# Methodology

Search Space Quality Score —— *如何评价hyperspace中的每一个search space的好坏?*

Latency-aware space quality of Q-T score.

Treat a space with good quality if its best searched subnets achieve optimal quantized accuracy under latency constraints T.

如果说一个search space 中最好的那个subnet在延迟限制下达到了最佳的量化精度，那么我们认为这样一个search space 是好的

用Q-T Score 来描述一个search space 的好坏

we treat every constraint equally important, and define Q-T score as the sum of each constraint:  $Q(\mathcal{A}, T_1, \dots, T_n) = Q(\mathcal{A}, T_1) + Q(\mathcal{A}, T_2) + \dots + Q(\mathcal{A}, T_n)$ ,  $Q(\mathcal{A}, T_i)$  is defined as:

$$Q(\mathcal{A}, T_i) = \mathbb{E}_{\alpha \in \mathcal{A}, LAT(\alpha) \leq T_i} [Acc_{int8}(\alpha)] \quad (1)$$

$\alpha$ 是搜索空间中最佳的subnet

$Acc_{int8}(\alpha)$ 是top-1 quantized accuracy evaluated on ImageNet validation set

$LAT(\alpha)$ 是量化后在目标设备上的延迟

Q-T Score就是在一定延迟限制下，获得最佳量化精度模型的数学期望，显然我们希望这个数学期望越大越好

# Methodology

## Elastic Stage and Problem Formulation

a chainstructured search space



Each stage defines a range of configurations  $c$  (e.g., kernel size, channel width, depth) for a specific block type  $b$ , and allows NAS to find the optimal architecture settings.

$$\mathcal{A} = STEM \circ E_{b,c}^1 \dots \circ E_{b,c}^N \circ HEAD \quad (2)$$

# Methodology

## Elastic Stage and Problem Formulation

### Problem definition

**Operator type  $b$**  and **configuration  $c$**  are two crucial objectives when searching quantization-friendly search space.

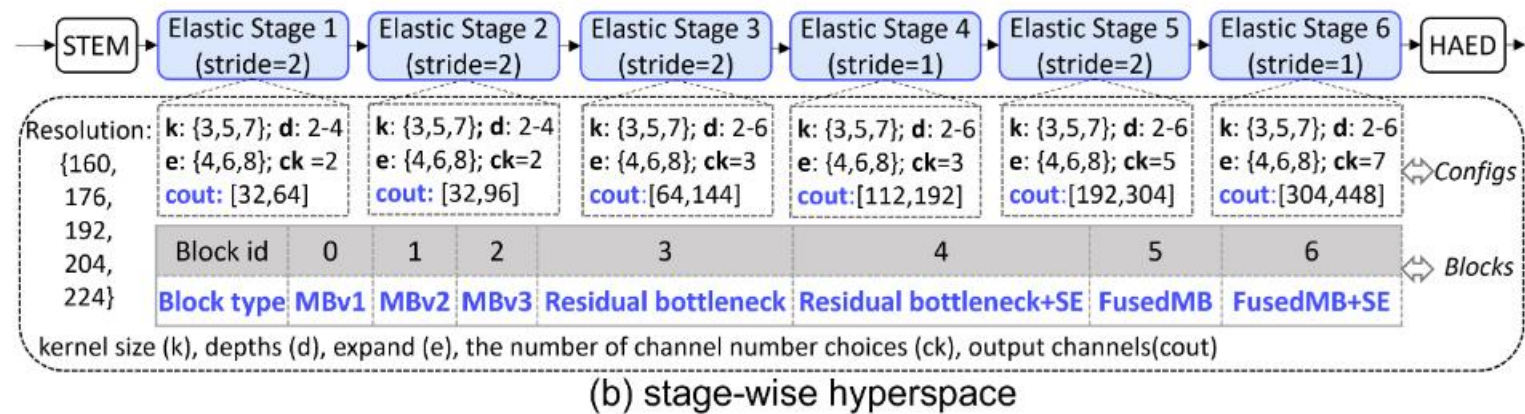
The task of space search can be simplified to find a search space **with optimal elastic stages**

$$\mathcal{A}(E_{b,c}^1 \circ E_{b,c}^2 \circ \dots \circ E_{b,c}^N)^* = \arg \max_{E_{b,c}^i \in \mathcal{H}^i} \mathcal{Q}(\mathcal{A}(E_{b,c}^1 \circ E_{b,c}^2 \circ \dots \circ E_{b,c}^N), T)$$

# Methodology

## Searching the Search Space

### Hyperspace design



Search by two dimensions

- Block type *b*
- Output channel width *cout*

A sampled search space is encoded by a sequential elastic stages.

Hyperspace has 10<sup>9</sup> candidate search spaces!



# Methodology

## Searching the Search Space

Evolutionary space search——aging evolution[1] 类似于遗传算法

1. 我们首先随机初始化一组  $P$  个搜索空间，构成一个种群，其中每个采样空间被编码为  $(E_{b,c}^1 \circ E_{b,c}^2 \circ \dots \circ E_{b,c}^N)$
2. 对每一个采样的搜索空间都使用 Q-T score快速评估
3. 接下来进入mutation iterations
4. 在每次迭代中，我们再从hyperspace中随机采样 $S$ 个搜索空间，并选择得分最高的候选者作为parent
5. mutation: 通过分别改变parent的block type 和 widths以生成两个child搜索空间。
6. 我们评估这两个搜索空间的Q-T分数，并将它们添加到当前种群中。
7. 最古老的 2个space 被删除，然后重复迭代。
8. 在所有迭代完成后，我们收集所有采样空间并选择Q-T分数最高的空间作为最终搜索空间。

# Evaluation

## Setup

- The INT8 latency constraints are {8, 10, 15, 20, 25} ms for Intel CPU and {15, 20, 25, 30, 35} ms for Pixel4.
- For each device, we search 5k search spaces in total and return the one with highest Q-T score.
- The population size  $P$  is 500 and sample size  $S$  is 125.

# Evaluation

Comparison with SOTA search spaces.

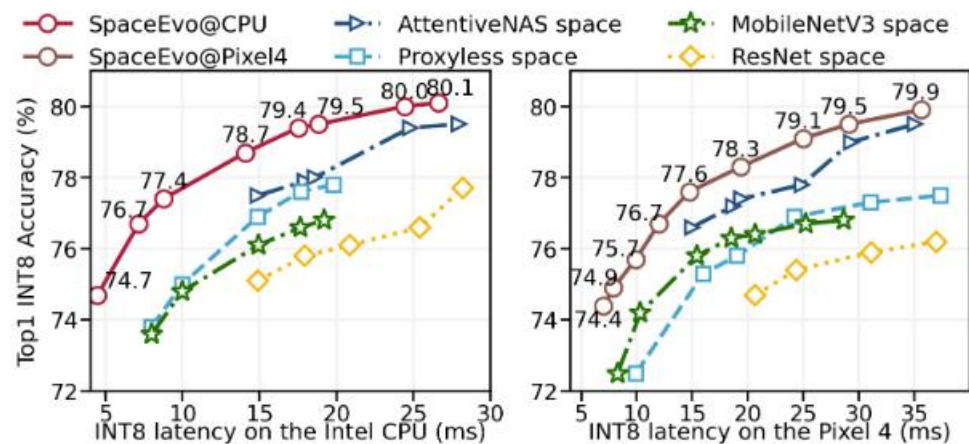


Figure 5. Best searched INT8 models with comparison to state-of-the-art NAS search spaces. Our searched spaces are proven to be the most quantization-friendly for the target device.

# Evaluation

## SpaceEvo under diverse latency constraints

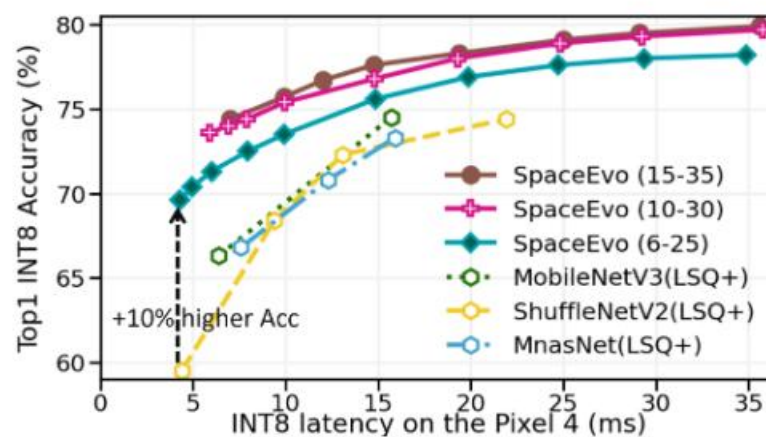


Figure 6. Search space design under diverse INT8 latency constraints. SpaceEvo (6-25 ms) delivers superior tiny INT8 models.

# The Effectiveness of Discovered INT8 Models

Table 1. ImageNet results compared with SOTA quantized models on two devices. \*: latency compared to FP32 inference.

| (a) Results on the Intel VNNI CPU with onnxruntime |              |                     |             |              |       |
|--|--------------|---------------------|-------------|--------------|-------|
| Model  | Acc%<br>INT8 | CPU Latency<br>INT8 | speedup*    | Acc%<br>FP32 | FLOPs |
| MobileNetV3Small                                   | 66.3         | 4.4 ms              | 1.1×        | 67.4         | 56M   |
| <b>SeqNet@cpu-A0</b>                               | <b>74.7</b>  | <b>4.4 ms</b>       | <b>2.0×</b> | 74.8         | 163M  |
| MobileNetV2  | 71.4         | 7.3 ms              | 2.2×        | 72.0         | 300M  |
| ProxylessNAS-R                                     | 74.6         | 8.8 ms              | 1.8×        | 74.6         | 320M  |
| OQAT-8bit  | 74.8         | 9.8 ms              | 1.8×        | 75.2         | 214M  |
| MobileNetV3Large                                   | 74.5         | 10.3 ms             | 1.5×        | 75.2         | 219M  |
| OFA (#25)  | 75.6         | 11.2 ms             | 1.5×        | 76.4         | 230M  |
| <b>SeqNet@cpu-A1</b>                               | <b>77.4</b>  | <b>8.8 ms</b>       | <b>2.4×</b> | 77.5         | 358M  |
| APQ-8bit   | 73.6         | 15.0 ms             | 1.5×        | 73.6         | 297M  |
| AttentiveNAS-A0                                    | 76.1         | 15.1 ms             | 1.4×        | 77.3         | 203M  |
| OQAT-8bit  | 76.3         | 14.9 ms             | 1.7×        | 76.7         | 316M  |
| EfficientNet-B0                                    | 76.7         | 18.1 ms             | 1.6×        | 77.3         | 390M  |
| <b>SeqNet@cpu-A2</b>                               | <b>78.5</b>  | <b>14.1 ms</b>      | <b>2.4×</b> | 78.8         | 638M  |
| APQ-8bit   | 74.9         | 20.0 ms             | 1.5×        | 75.0         | 393M  |
| OQAT-8bit  | 76.9         | 19.5 ms             | 1.6×        | 77.3         | 405M  |
| AttentiveNAS-A1                                    | 77.2         | 22.4 ms             | 1.4×        | 78.4         | 279M  |
| AttentiveNAS-A2                                    | 77.5         | 22.5 ms             | 1.3×        | 78.8         | 317M  |
| <b>SeqNet@cpu-A3</b>                               | <b>79.5</b>  | <b>18.9 ms</b>      | <b>2.6×</b> | 79.6         | 981M  |
| FBNetV2-L1   | 75.8         | 25.0 ms             | 1.2×        | 77.2         | 325M  |
| FBNetV3-A  | 78.2         | 27.7 ms             | 1.3×        | 79.1         | 357M  |
| <b>SeqNet@cpu-A4</b>                               | <b>80.0</b>  | <b>24.4 ms</b>      | <b>2.4×</b> | 80.1         | 1267M |

| (b) Results on the Google Pixel 4 with TFLite |              |                        |             |              |       |
|---|--------------|------------------------|-------------|--------------|-------|
| Model   | Acc%<br>INT8 | Pixel4 Latency<br>INT8 | speedup*    | Acc%<br>FP32 | FLOPs |
| MobileNetV3Small                              | 66.3         | 6.4 ms                 | 1.3×        | 67.4         | 56M   |
| <b>SeqNet@pixel4-A0</b>                       | <b>73.6</b>  | <b>5.9 ms</b>          | <b>2.1×</b> | 73.7         | 107M  |
| MobileNetV2                                   | 71.4         | 16.5 ms                | 1.9×        | 72.0         | 300M  |
| ProxylessNAS-R                                | 74.6         | 18.4 ms                | 1.8×        | 74.6         | 320M  |
| MobileNetV3Large                              | 74.5         | 15.7 ms                | 1.5×        | 75.2         | 219M  |
| APQ-8bit                                      | 74.6         | 14.9 ms                | 2.0×        | 74.4         | 340M  |
| OFA (#25)                                     | 75.6         | 14.8 ms                | 1.7×        | 76.4         | 230M  |
| OQAT-8bit                                     | 75.8         | 15.2 ms                | 1.9×        | 76.2         | 287M  |
| AttentiveNAS-A0                               | 76.1         | 15.2 ms                | 2.0×        | 77.3         | 203M  |
| <b>SeqNet@pixel4-A1</b>                       | <b>77.6</b>  | <b>14.7 ms</b>         | <b>2.2×</b> | 77.7         | 274M  |
| APQ-8bit                                      | 75.1         | 20.0 ms                | 1.9×        | 75.1         | 398M  |
| OQAT-8bit                                     | 76.5         | 20.4 ms                | 1.8×        | 76.8         | 347M  |
| AttentiveNAS-A1                               | 77.2         | 21.1 ms                | 2.0×        | 78.4         | 279M  |
| AttentiveNAS-A2                               | 77.5         | 22.7 ms                | 2.0×        | 78.8         | 317M  |
| <b>SeqNet@pixel4-A2</b>                       | <b>78.3</b>  | <b>19.4 ms</b>         | <b>2.3×</b> | 78.4         | 402M  |
| FBNetV2-L1                                    | 75.8         | 26.7 ms                | 1.5×        | 77.2         | 325M  |
| OQAT-8bit                                     | 77.0         | 29.9 ms                | 1.7×        | 77.2         | 443M  |
| FBNetV3-A                                     | 78.2         | 30.5 ms                | 1.5×        | 79.1         | 357M  |
| <b>SeqNet@pixel4-A3</b>                       | <b>79.5</b>  | <b>30.8 ms</b>         | <b>2.1×</b> | 79.5         | 591M  |
| EfficientNet-B0                               | 76.7         | 36.4 ms                | 1.7×        | 77.3         | 390M  |
| <b>SeqNet@pixel4-A4</b>                       | <b>79.9</b>  | <b>35.5 ms</b>         | <b>2.2×</b> | 80.0         | 738M  |

# Thoughts

## Strength

- automatically design a quantization-friendly space for target device, which delivers superior INT8 quantized models with SOTA efficiency on real-world edge devices.
- the proposed quality score (Q-T score) is a good approximate metric of the search space quality.
- Its potential extends to applications for other crucial deployment metrics, such as energy and memory consumption, enhancing the sustainability of edge computing solutions.

## Weakness

- Experiments are mostly done with two mobile CPUs. It is uncertain whether the on-device quantization efficiency analysis in this paper still holds for mobile GPUs and even mobile NPUs or not.
- The Experiments devices are too few, which can not make the results or observation convincing.



***Thank You !***

*Ye Wan 2023.11.20*